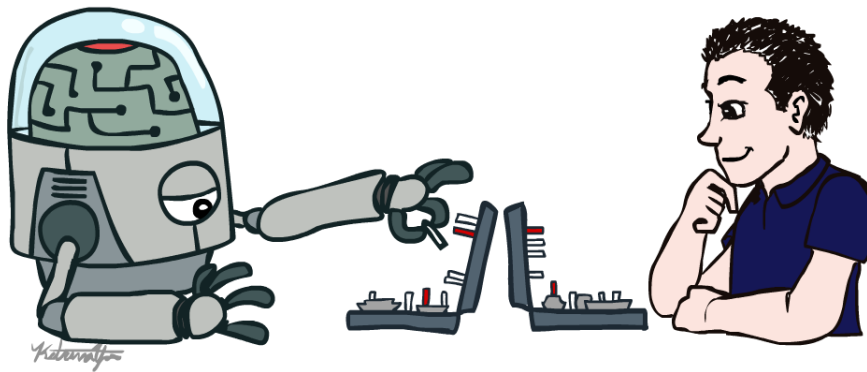


CSE 3521: Introduction to Artificial Intelligence



[These slides are partially adapted from the [UC Berkeley. CS188 Intro to AI](#) at UC Berkeley]

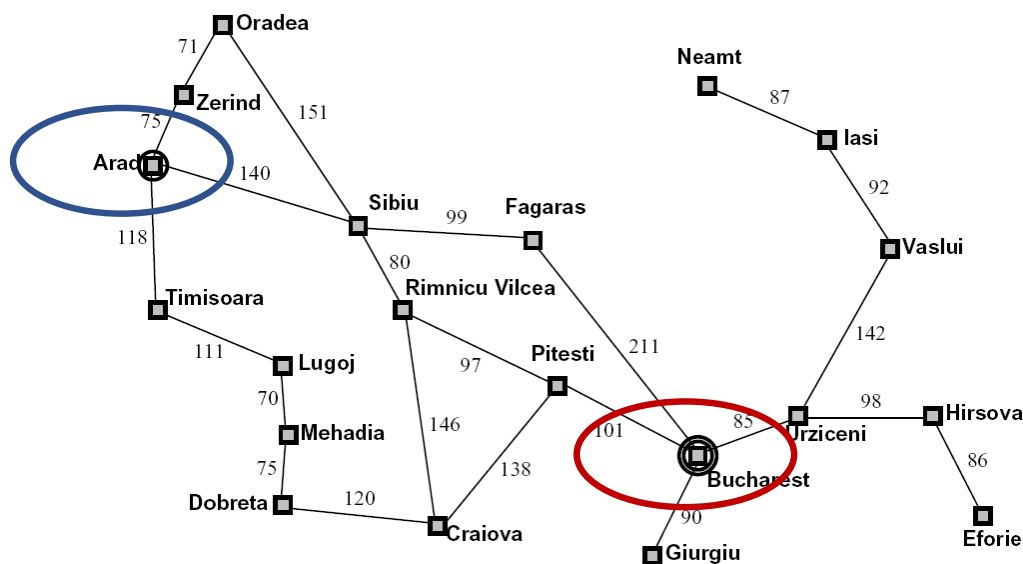


THE OHIO STATE UNIVERSITY

Search Problems


- A **search problem** consists of:
 - A state space
 - A successor function
(with **actions**, **costs**)
 - A start state and a goal test
- A **solution** is a sequence of actions (a plan) which transforms the start state to a goal state

Example: Traveling in Romania



- State space:
 - Cities
- Successor function:
 - Roads: Go to adjacent city with cost = distance
- Start state:
 - Arad
- Goal test:
 - Is state == Bucharest?
- Solution?
 - If there exists one, it is a sequence of cities from Arad to Bucharest


Searching with a Search Tree



- Search:

- Expand out potential plans (tree nodes)
- Maintain a **fringe** of partial plans under consideration
- Try to expand as few tree nodes as possible


Searching with a Search Tree



- Search:

- Expand out potential plans (tree nodes)
- Maintain a **fringe** of partial plans under consideration
- Try to expand as few tree nodes as possible

Searching with a Search Tree



- Search:
 - Expand out potential plans (tree nodes)
 - Maintain a **fringe** of partial plans under consideration
 - Try to expand as few tree nodes as possible

General Tree Search

```
function TREE-SEARCH( problem, strategy ) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

- Important components: (1) Fringe (2) Expansion (3) Exploration strategy
- Main question: which fringe nodes to explore?

Uninformed vs. Informed Search

- Uninformed search
 - Given no information about problem (other than its definition)
 - Find solutions to problems by systematically generating new states and testing for goal
- Informed search
 - Given some ideas of where to look for solutions
 - Use problem-specific knowledge

Uninformed vs. Informed Search


- Uninformed search

- Given no information about problem (other than its definition)
- Find solutions to problems by systematically generating new states and testing for goal

- Informed search

- Given some ideas of where to look for solutions
- Use problem-specific knowledge



Example: Tree Search



Depth-First Search

Strategy: expand a
deepest node first

Implementation:
Fringe is a LIFO stack



Search Algorithm Properties


- **Complete:** Guaranteed to find a solution if one exists?
- **Optimal:** Guaranteed to find the least cost path?

- Search tree:

- b is the branching factor
- m is the maximum depth
- solutions at various depths


- Number of nodes in entire tree?

- $1 + b + b^2 + \dots + b^m = O(b^{m+1})$



Depth-First Search (DFS) Properties



- What nodes DFS expand?
 - Some left prefix of the tree.
 - Could process the whole tree!
 - If m is finite, takes time $O(b^m)$
- How much space does the fringe take?
 - Only has siblings on the current path to root, so $O(bm)$
- Is it complete?
 - m could be infinite, so only if we prevent cycles
- Is it optimal?
 - No, it finds the “leftmost” solution, regardless of depth or cost



Breadth-First Search


Strategy: expand a shallowest node first

Implementation: Fringe is a FIFO queue




Breadth-First Search (BFS) Properties

- What nodes does BFS expand?
 - Processes all nodes above shallowest solution
 - Let depth of shallowest solution be s
 - Search takes time $O(b^s)$
- How much space does the fringe take?
 - Has roughly the last tier, so $O(b^s)$
- Is it complete?
 - s must be finite if a solution exists, so yes!
- Is it optimal?
 - Only if costs are all 1 (more on costs later)




Iterative Deepening Search

- Idea: get DFS's space advantage with BFS's time / shallow-solution advantages
 - Run a DFS with depth limit 1. If no solution...
 - Run a DFS with depth limit 2. If no solution...
 - Run a DFS with depth limit 3.
- Isn't that wastefully redundant?
 - Generally most work happens in the lowest level searched, so not so bad!
- A Preferred method with large search space and depth of solution not known



Cost-Sensitive Search





BFS finds the shortest path in terms of number of actions.
It does not find the least-cost path. We will now cover
a similar algorithm which does find the least-cost path.

Uniform Cost Search (USC)


Strategy: expand a cheapest node first:

Fringe is a priority queue
(priority: cumulative cost)



Uniform Cost Search (UCS) Properties

- What nodes does UCS expand?
 - Processes all nodes with cost less than cheapest solution!
 - If that solution costs C^* and arcs cost at least ε , then the “effective depth” is roughly C^*/ε
 - Takes time $O(b^{C^*/\varepsilon})$ (exponential in effective depth)
 - Additional cost of priority queue
- How much space does the fringe take?
 - Has roughly the last tier, so $O(b^{C^*/\varepsilon})$
- Is it complete?
 - Assuming best solution has a finite cost and minimum arc cost is positive, yes!
- Is it optimal?
 - Yes!



The One Queue

- All the search algorithms are the same except for fringe/exploration strategies
 - Conceptually, all fringes are priority queues (i.e. collections of nodes with attached priorities)
 - Practically, for DFS and BFS, you can avoid the $\log(n)$ overhead from an actual priority queue, by using stacks and queues
 - Can even code one implementation that takes a variable queuing object

